

Intro

- Jan Koteck
 - jan@koteck.net
 - @JanKoteck
 - MapDB author (database engine)
 - Java developer for past 12 years
 - Last 3 years independent consultant for MapDB
- Slides and code at
 - <https://github.com/jankoteck/talk-save-java-memory>

Intro

- This talk is about:
 - Reducing memory footprint
 - Object overhead
 - Primitive variables and arrays
 - Alternative Java Collections
- Not about:
 - Garbage Collection
 - Off-heap
 - Performance
 - Memory leaks

Measure memory consumption

- `Runtime.getRuntime().freeMemory()` difference
- Allocate until `OutOfMemoryError`
- Profiler
- Class or heap dump analyze such as Java Object Layout

Runtime.getRuntime().freeMemory()

- Get allocated memory stats
- Create new objects
- Call GC many times
- Get allocated memory difference
- Not very precise, but easy
- May not always work, some JVMs ignore System.gc()

```
NanoBench nanoBench = NanoBench.create();
nanoBench.memoryOnly().warmUps(2).measurements(10)
.measure(String.format("Measure memory for %d keys", i),
    new Runnable() {
        @Override public void run() {
            Random setRandom = new Random(4532);
            ref = new HashSet<Integer>(cnt);
            while (ref.size() < cnt) {
                ref.add(setRandom.nextInt());
            }
        }
    });
nanoBench.getMemoryBytes()
```


OutOfMemoryError















- Memory limit (-Xmx1G)
- Add entries to array in loop
- Print number of items
- JVM eventually crashes

```
try{
    //add new entry in loop
    while(true)
        array[arrayPos++] = new HashMap();
}catch(OutOfMemoryError e){
    memoryConsumption =
        Runtime.getRuntime().maxMemory()
            / counter;

    return;
}
```

Profiler

- Shows memory consumed by single class
 - (byte[], Object[])
- VisualVM is easy to use
- Flight Recorder is best
 - Free outside production env.

Class Name - Allocated Objects	Bytes Allocated	Bytes Allocated	Objects Allocated
char[]		50,706,872 (23.1%)	290,474 (8.2%)
int[]		28,125,808 (12.8%)	69,664 (2%)
byte[]		21,194,384 (9.6%)	98,475 (2.8%)
java.lang.Object[]		12,823,272 (5.8%)	277,607 (7.8%)
java.lang.String		6,310,536 (2.9%)	262,939 (7.4%)
java.lang.Class		3,817,200 (1.7%)	34,887 (1%)
com.intellij.openapi.vfs.newvfs.impl.Virtual...		3,748,640 (1.7%)	117,145 (3.3%)
java.util.TreeMap\$Entry		3,563,840 (1.6%)	89,096 (2.5%)
java.io.ObjectStreamClass\$WeakClass...		2,317,184 (1.1%)	72,412 (2%)
long[]		2,252,032 (1%)	7,631 (0.2%)
com.intellij.psi.impl.source.tree.java.PsiJa...		2,044,800 (0.9%)	51,120 (1.4%)
com.intellij.util.containers.IntObjectLinke...		1,921,504 (0.9%)	60,047 (1.7%)
com.intellij.psi.impl.source.tree.PsiWhite...		1,842,240 (0.8%)	46,056 (1.3%)
com.intellij.util.text.ByteArrayCharSeque...		1,776,648 (0.8%)	74,027 (2.1%)

Java Object Layout

- Command line tool to analyze memory layout (and space usage)
- Most advanced
- <http://openjdk.java.net/projects/code-tools/jol/>
- Written by Aleksey Shipilev
- Example:

```
$ java -jar jol-cli/target/jol-cli.jar internals java.util.HashMap
```

```
Instance size: 48 bytes (reported by Instrumentation API)
```

```
Space losses: 0 bytes internal + 4 bytes external = 4 bytes total
```

Instance overhead in Java

- Pointer overhead
 - 4 to 8 bytes
 - depends on JVM and amount of memory
 - 64bit JVM 8 uses compressed pointer when possible
- Class info and object header
 - 12 bytes
- Object alignment padding
 - Between 0 and 15

35 bytes worst case scenario

java.lang.Integer overhead

- One object + one reference
- Total overhead 16 bytes + reference (4 to 8 bytes)

JOL output for java.lang.Integer

OFFSET	SIZE	TYPE	DESC
0	12		(object header)
12	4	int	Integer.value

total size: **16 bytes**

java.lang.String overhead

- Two objects and two references
- One int field (String.hash)
- Two bytes per char (UTF16)
- Other problems:
 - Weak hash (many collisions)
 - Defensive char[] copies
 - Substring create copies etc..

JOL output for java.lang.String

OFFSET	SIZE	TYPE	DESC
0	4	(object header)	
4	4	(object header)	
8	4	(object header)	
12	4	char[]	String.value
16	4	int	String.hash
20	4	(alignment loss)	

total size: **24 bytes**

String.intern()

- Reduces number of instances
- Each string creates new weak ref (huge GC overhead)
- Long.valueOf(11) is similar

```
name.toLowerCase().intern()  
    ==  
name.toLowerCase().intern()
```

Represent String as char[]

- Field is char[] or byte[]
- Get/set takes string
- GC overhead, but short lived objects are fast to collect
- Reduced memory with many objects (large Maps etc)

```
public class Person{  
  
    private char[] name;  
  
    public String getName(){  
        return new String(name);  
    }  
  
    public void setName(String name){  
        this.name = name.toCharArray();  
    }  
}
```


char[] as key in Map

- char[] has no equals() and hashCode() methods
- TreeMap key can use Comparator<char[]>
 - com.google.common.primitives.Chars.lexicographicalComparator()
- java.util.HashMap can not change hash strategy
 - Object2ObjectOpenCustomHashMap from FastUtil has Hash Strategies
 - String.hashCode() is broken, many collisions
 - Part of spec, can not be changed :-)
 - ConcurrentHashMap Java8 has sorted lists to handle hash duplicates :-)

char[] as value in Map

- Works (mostly)
- Concurrent Compare-And-Swap operations broken
 - char[].equals() uses identity
 - ConcurrentMap.replace(K key, V oldValue, V newValue);
 - MapDB Maps allow custom equality

Bitwise

```
class Address{  
    String street;  
    int houseNumber;  
}
```

- Replace street with number from dictionary
- Combine streetId and houseNumber into single long number
 - 5 bytes for streetId, 3 bytes for house number and flags
- Address become primitive number, no objects or references

Primitive key / value in collections

- `HashMap<Long,Long>` uses `Object[]` for keys and values
- ~40 bytes overhead for each key-value pair
- Use `long[]` or `int[]` instead
- Data locality → CPU cache friendly
 - Binary search much faster
 - Faster fetch of values
- Almost zero Garbage Collection
 - entire `HashMap` uses two objects
 - No objects allocated for new key or value

FastUtil

- <http://fastutil.di.unimi.it/>
- Autogenerated classes with primitive keys and values
- Int2IntMap, Int2BooleanMap, Int2ObjectMap<V>
- Many classes, 17MB jar

FastUtil

```
// create new map
Int2ObjectMap<String> map =
    new Int2ObjectOpenHashMap();

// put/get has primitive arg, no unboxing!
map.put(1, "aa");
map.get(1);

// also implements full mapdb interface
Map<Integer, String> map2 = map;
```

- Parallel interface hierarchy
- Get / set has primitive arg version, no unboxing!
- Implements Map interface

Chronicle Map

- ConcurrentHashMap
- Keys and values are serialized outside heap
- Supports any serializable data type
- No objects or references on heap, memory efficient
- <http://chronicle.software/products/chronicle-map/>

```
Map<Long, long[]> graph =  
ChronicleMap  
    .of(Long.class, long[].class)  
    .entries(1_000_000_000L)  
    .averageValue(new long[150])  
    .create();
```

Conclusion

- Avoid too many object instances in long lived objects
- Short lived objects are fine, when traded for smaller memory
- Email: jan@kotek.net
- Slides and code at
 - <https://github.com/jankotek/talk-save-java-memory>